Pedro Goulart Rodrigues
Senior Data Platform Engineer | Data Platforms & Real-time Streaming
Porto Alegre, Brazil · hi@pedrorod.com
linkedin.com/in/pedro-goulart-rodrigues

Profile

Senior Data Engineer with 7+ years in data, the last four focused on cloud. I build and own batch and streaming platforms end to end — architecture, reliability and cost — at large scale: 1,000+ Airflow DAGs, 300+ Kafka streaming pipelines and tens of terabytes a day for 1,000+ internal users. I was part of the AWS-to-GCP migration and led the SAP streaming ingestion. I also work at the intersection of data and AI, building an AI-based solution to detect sensitive data and supporting LLM/agent platform initiatives. What I enjoy most is building scalable systems and data products that help teams move faster with data they can trust.

Skills

Core: Python, SQL, Spark, Kafka, Airflow, dbt, DuckDB

Cloud: GCP (BigQuery, Dataproc, Cloud Storage, Composer), AWS (Redshift, S3, Athena)

Platform & DevOps: Kubernetes, Terraform, Helm, GitHub Actions, CI/CD, FinOps

Data: Data Warehousing, Streaming, ETL/ELT, Data Governance, SLA Monitoring, Data Products, APIs

Tools: Power BI, Tableau, SSIS, Pentaho

Languages: Portuguese (Native), English (Full Professional)

Professional Experience

Arco Educação — Senior Data Platform Engineer

September 2026 – Present | São Paulo, Brazil

Building scalable batch and streaming data platforms and data products for one of the largest education groups in Brazil.

Grupo SBF — Data Platform Engineer

March 2022 – September 2026 | São Paulo, Brazil

- Built and maintained a large-scale data platform on GCP supporting 1,000+ Airflow DAGs and 1,000+ internal users across analytics and reporting.
- Designed and operated 300+ real-time streaming pipelines with Kafka, processing tens of TB daily across domains.
- Led the SAP data streaming integration with Kafka and SAP Datasphere, including schema-change detection and proactive alerts for breaking changes.
- Took part in the AWS-to-GCP migration, redesigning pipelines and decommissioning legacy systems.
- Reduced Kubernetes costs by ~23% and storage costs by 15% through infrastructure optimization, without touching SLAs.
- Designed task-level SLA monitoring for Airflow and built dashboards to track adherence and pipeline performance.
- Built an AI-based solution to detect sensitive data across datasets, improving security and governance.
- Managed API integrations for third-party ingestion and built reusable batch and streaming ingestion frameworks.

Tech Stack: GCP · BigQuery · Dataproc · Cloud Storage · Airflow · Kafka · Kubernetes · Terraform · Helm · CI/CD · Python · SQL · DuckDB · SAP (AecorSoft, Datasphere) · Data Governance

RCI Analytics Consulting — Data Engineer
October 2019 – March 2022 | Porto Alegre, Brazil

- Built end-to-end data pipelines integrating 10+ data sources and processing ~5TB daily.
- Reduced pipeline failures by 30% by optimizing SQL queries and API scripts.
- Automated ingestion from APIs, FTPs and cloud storage, and modeled star and snowflake warehouses.
- Delivered dashboards used by 300+ users for operational and executive decisions.
- Worked with retail, finance and logistics clients, defining KPIs and reporting needs.

Tech Stack: Python · SQL · SSIS · Pentaho · Tableau Prep · Power BI · Tableau · Redshift · S3 · Oracle · MSSQL · PostgreSQL

Grupo Oleoplan — Data
January 2019 – October 2019 | Porto Alegre, Brazil

- Audited invoicing and sales data to support tax reporting and compliance.
- Built Power BI dashboards to track financial KPIs and reduce manual reporting.
- Provided IT support and system routines for finance and operations, plus data cleaning and validation.

Stack: Power BI · Excel · Data Auditing · Reporting · Internal Systems

Education

Descomplica Faculdade Digital — System Analysis
July 2023 – December 2025

UniRitter — Systems Analysis
2019 – 2022